

Record I/O

CS3411 Fall 2021

Program Two

Due: October 15, 2021, 8am

Unix abstraction of files as byte streams is powerful and versatile. However, there are times having a file implementation which is aware of *records* would simplify developing certain types of programs. A record based file i/o will operate on units of records which can be fixed or varying length. Although the Unix abstraction of byte streams work well with fixed length records, dealing with variable length records is notably more difficult. A typical example would be the need to read and write a given file line at a time where the length of each line varies.

In this assignment, our goal is to create a set of abstractions imitating regular open/read/write/lseek/close kernel calls but implement record based I/O semantics. The abstractions and their parameters are listed later in this description.

Representing Record Files

The easiest way to represent record files would be to encode the length of each record in the file. However, this approach would make the file contents inaccessible to other Unix utilities. Instead, we will create an *index* file which accompanies the data file and store the record length and position information in the index file. Each record in the data file will then be represented using a *record descriptor* with the following layout:

```
struct record_descriptor
{
    int position; //byte offset relative to the beginning of the data file
    int length;   //length of record in bytes
};
```

The index file is a binary file that is read and written by performing a kernel i/o for the size of the structure:

```
struct record_descriptor r;

r.position = record_position;
r.length = strlen(buf);

n = write(inx, &r, sizeof(struct record_descriptor));
```

The index file for record i/o is named by following the convention `.rinx.<data file name>` (i.e., if the data file is named `myfile.txt`, its index will be `.rinx.myfile.txt`). This convention conveniently hides the index files while permitting easy access to them based on the data file name.

As an example, assume the following program fragment is executed using the abstractions you have created:

```
char * record;

if (( d = rio_open("myfile.txt", O_RDWR | O_CREAT, 0600)) < 0) all_hands_abandon_ship(RIO_OPEN);
if (( n = rio_write(d,"Systems\n",8)) != 8) all_hands_abandon_ship(RIO_ERROR_WRITE);
if (( n = rio_write(d,"programming is cool\n",20)) != 20) all_hands_abandon_ship(RIO_ERROR_WRITE);
if (rio_close(d) < 0) all_hands_abandon_ship(RIO_CLOSE);

if (( n = rio_lseek(d,1, SEEK_SET)) != 1) all_hands_abandon_ship(RIO_ERROR_LSEEK);

record = rio_read(d, &n);
if (n == 0) { fprintf(stderr,"Unexpected EOF on record i/o\n"); exit(-1);}
if (n < 0) all_hands_abandon_ship(RIO_ERROR_READ);
```

```
write(1,record, n);
free(record);
```

If there are no errors, this program should print *programming is cool* as `rio_lseek` seeks the file to the beginning of the second record in the file (`rio_lseek(d,1, SEEK_SET)` is zero relative). The content of the resulting file is shown below. Note the index file contents are shown in decimal and comma separated for readability:

Data File	Index file
-----	-----
Systems	0, 8
programming is cool.	8, 20

Simplifications and Notes

In order to complete the project at a reasonable time, you can make the simplification that if the file is not at the end of file and a write is performed, the write length should be less than or equal to the length of the existing record in the file.

Your developed abstractions should permit multiple record files to be opened and manipulated.

Semantics of Record i/o Routines

`rio_open`

Prototype: `int rio_open(const char *pathname, int flags, mode_t mode)`

Open the data file and index file. If create is requested, open both the data file and the index file with `O_CREAT`. If `O_RDONLY` or `O_RDWR` is requested, make sure that the index file is present and return an error otherwise. On success the data file's descriptor should be returned and the file descriptor for the index file must be maintained within the abstraction.

`rio_read`

Prototype: `void * rio_read(int fd, int * return_value)`

Allocate a buffer large enough to hold the requested record, read the next record into the buffer and return the pointer to the allocated area. The I/O result should be returned through the `return_value` argument. On success this will be the number of bytes read into the allocated buffer. If any system call returns an error, this should be communicated to the caller through `return_value`.

`rio_write`

Prototype: `int rio_write(int fd, const void *buf, int count);`

Write a new record. If appending to the file, create a record descriptor and fill-in the values. Write the descriptor to the index file and the supplied data to the data file for the requested length. If updating an existing record, read the record descriptor, check to see if the new record fits in the allocated area and rewrite. Return an error otherwise.

`rio_lseek`

Prototype: `int rio_lseek(int fd, int offset, int whence);`

Seek both files (data and index files) to the beginning of the requested record so that the next I/O is performed at the requested position. The offset argument is in terms of records not bytes (relative to whence). *whence* assumes the same values as `lseek whence` argument.

rio_close

Prototype: `int rio_close(int fd);`

Close both files. Even though a single descriptor is passed along, your abstraction must close the other file as well. It is suggested that you return the descriptor obtained by opening the data file to the user and keep the index file descriptor number in the abstraction and associate them.

Pragmatics

Source Template

You are provided a source template that includes:

- Makefile with the following commands:
 - all - compile project into *recordio*, *testio*, *indexer*.
 - clean - removes executable and object files.
 - submission - generates prog2.tgz with all source files needed for submission. Upload this file to Canvas.
- recordio.h - header with prototypes for recordio functions
- recordio.c - C source where you will implement all recordio calls.
- recordTests.c - C source where you can implement tests that call the functions defined in recordio.
- indexer.c - Source where you will implement the standalone program indexer.
- testio.c - Source where you will implement the utility program testio.

Submission

Your submission must be written in C.

1. Develop the abstraction into a single C file which is titled *recordio.c* and a prototype file *recordio.h*. A user who wishes to use your abstraction should be able to do so by including *recordio.h* in their file and linking with *recordio.o*.
2. Develop a single standalone program called *indexer* which creates an index file for a text file whose title is supplied as an argument to the program. This program should read the text file from beginning till end, find the beginning and ending of each line (ending with the newline character), create a descriptor and write the descriptor to the created index file. The program should not modify the text file which is supplied as an argument.
3. Develop a test program called *testio* which includes *recordio.h* and is linked against *recordio.o*. This program should expect a single argument which supplies a file name. The program should *rio_open* the supplied argument and report any errors to the user. Upon a successful open it should execute a series of *rio_read* statements, read the file one record at a time and write each record to the standard output as shown above.
4. Test your *testio.c* and *indexer*. Use *indexer* to index *indexer.c* and *testio* to display the file line at a time.

Use Canvas to submit a tar file named **prog2.tgz** generated by running "make submission". Your source modifications must include comments.